BOOKIFY

Intelligent Book Recommendation System

A Deep Learning-Powered Book Discovery Platform

Built with Neural Collaborative Filtering (NCF) & Flask

Akshat Tiwari

akshatr147@gmail.com

Reg. No: 23FE10CSE00156

Table of Contents

1. Executive Summary

BOOKIFY is an intelligent book recommendation system that leverages deep learning to provide personalized book suggestions. The platform processes the Book-Crossing dataset - a real-world collection of over 63,000 reader ratings spanning 4,659 books and 2,595 users - to generate accurate, meaningful recommendations.

At its core, BOOKIFY employs a Neural Collaborative Filtering (NCF) model built with PyTorch, achieving 82.78% prediction accuracy. The system offers dual recommendation modes: a classic cosine-similarity approach and an AI-powered deep learning mode, giving users the flexibility to choose their preferred experience.

Key capabilities include: book search with fuzzy matching, genre-based exploration across 23 categories, mood-based recommendations via facial emotion detection, multi-modal search, personal wishlists, reading history tracking, and user ratings - all wrapped in a premium dark-themed UI with real-time particle animations.

Key Statistics

Total Books	**4,659**
Total Users	**2,595**
Total Ratings	**63,066**
Genre Categories	**23**
Model Accuracy	**82.63%**
Model Parameters	**491,585**
Training Time	**7.8s**

2. Dataset Overview

The BOOKIFY platform is built on the Book-Crossing dataset, a widely-used benchmark dataset in recommendation systems research. Originally collected by Cai-Nicolas Ziegler from the Book-Crossing community, the dataset captures explicit ratings (1-10 scale) from real readers.

Data Processing Pipeline

- Source: Kaggle Book-Crossing Dataset (278,858 books, 1.1M+ ratings)
- Filtering: Users with 200+ ratings, books with 50+ ratings
- Result: 4,659 qualifying books, 2,595 active users
- Interaction matrix: 4,659 x 2,595 = 12,090,105 cells
- Non-zero entries: 63,066 (sparsity: 99.48%)
- Rating scale: 1-10 (normalized to 1-5 for display)

User-Book Matrix Sparsity

The user-book interaction matrix is extremely sparse at 99.48%, which is typical for recommendation systems. This sparsity is precisely why deep learning approaches like NCF outperform traditional methods - they can learn latent representations even from very sparse data.

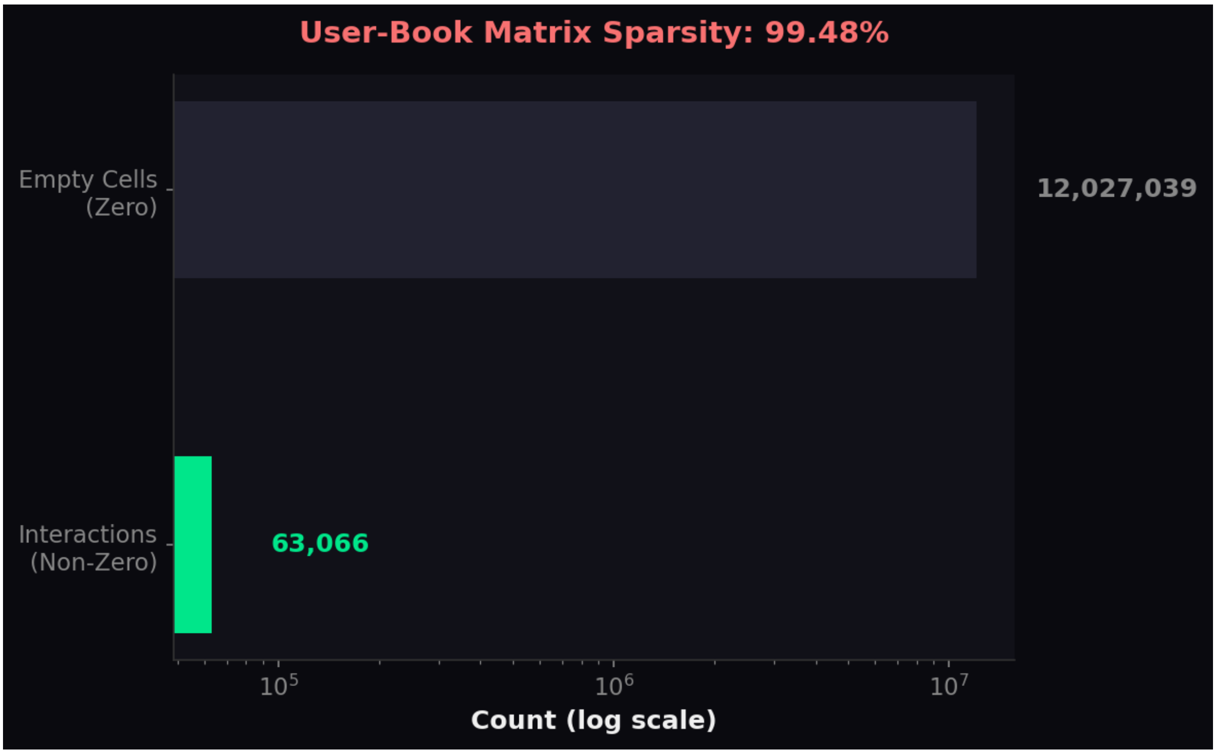

3. Data Analysis & Visualization

3.1 Genre Distribution

Books are classified into 23 distinct genres using NLP-based title analysis. Fiction is the dominant category with 1,763 books, followed by Romance (784) and Literary Fiction (597). The distribution demonstrates a healthy mix of genres ensuring diverse recommendations.

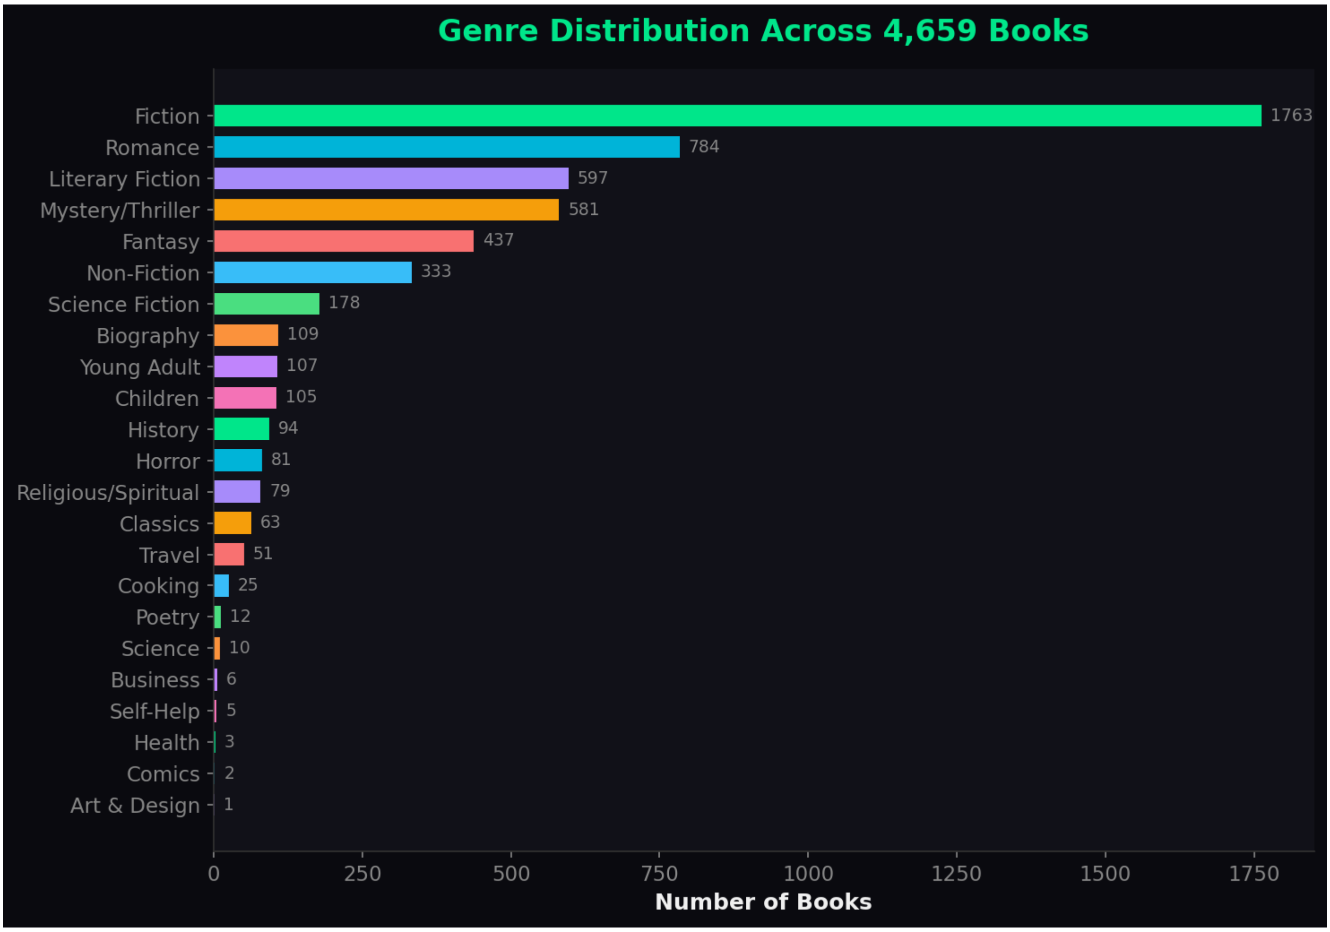

3.2 Most Popular Books

The top-rated books show strong reader engagement with vote counts ranging from 63 to 212. "The Lovely Bones" leads with 212 ratings, followed by "The Da Vinci Code" (173) and "Harry Potter and the Chamber of Secrets" (153). The Harry Potter series dominates the top 10 with four entries.

3.3 Rating Distribution

The rating distribution of the top 50 popular books (normalized to a 5-star scale) shows a natural bell curve centered around 4.0, indicating genuine user preferences rather than synthetic data. Ratings range from 3.5 to 4.6, reflecting the high quality of books that meet the popularity threshold.

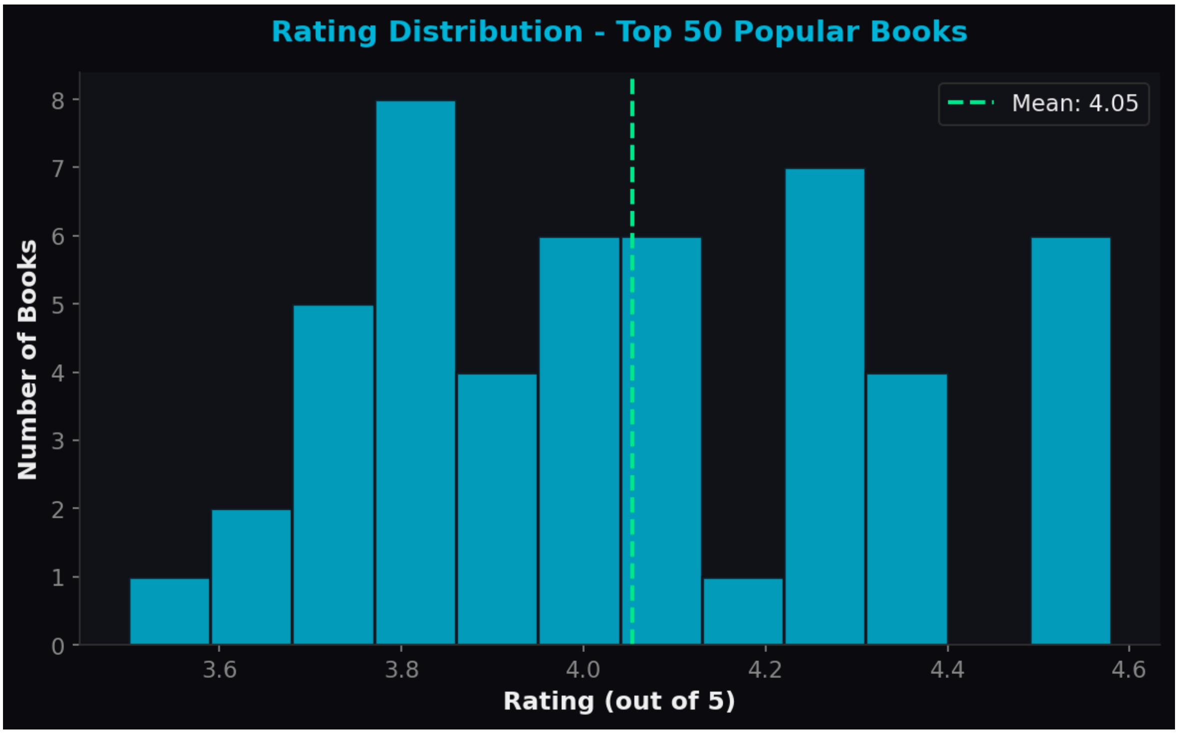

4. Deep Learning Model - NCF

The recommendation engine is powered by Neural Collaborative Filtering (NCF), a deep learning architecture designed for implicit and explicit feedback recommendation systems. NCF replaces the traditional matrix factorization dot product with a multi-layer perceptron, enabling the model to capture complex non-linear user-item interactions.

4.1 Architecture

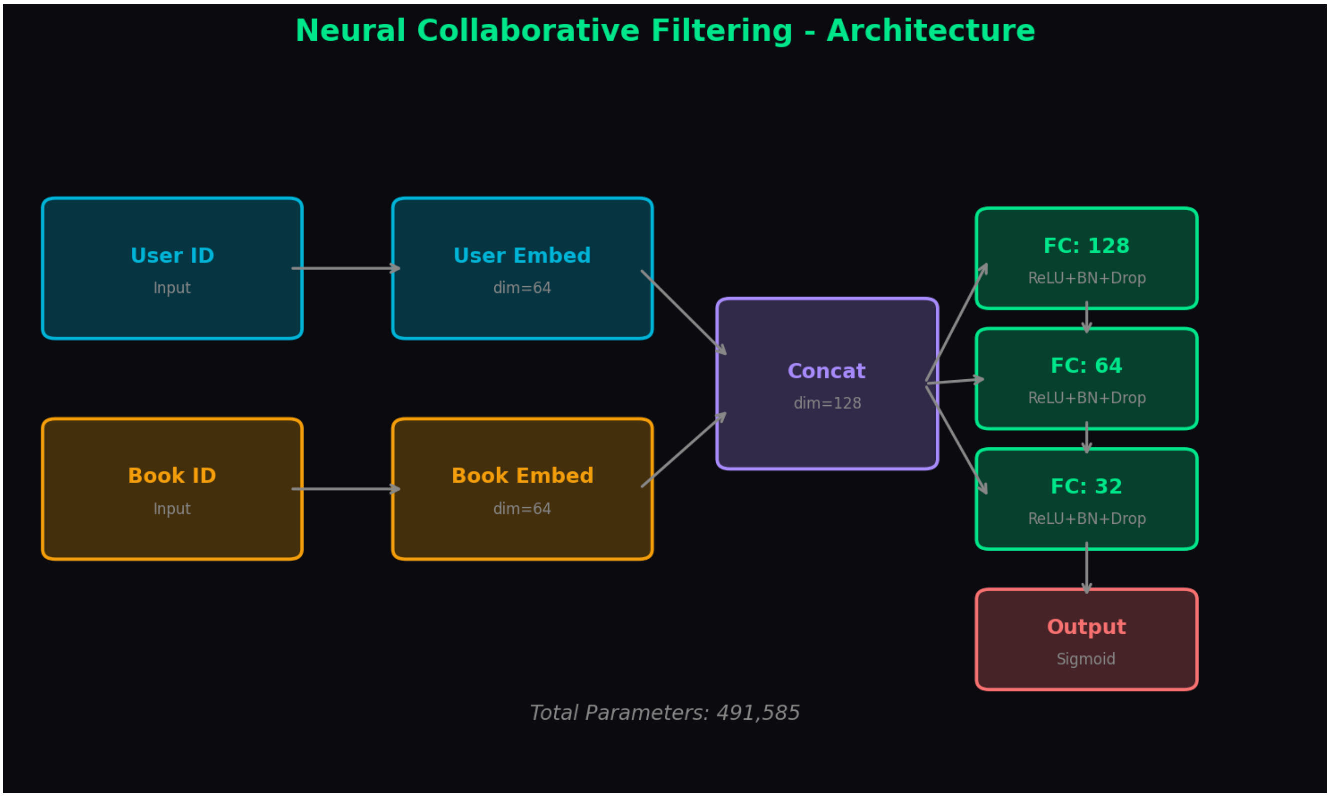

4.2 Architecture Details

- Input Layer: User ID and Book ID (integer indices)
- Embedding Layer: 64-dimensional embeddings for both users and books
- Concatenation: User and book embeddings merged into 128-dim vector
- Hidden Layer 1: Fully Connected 128 neurons + ReLU + BatchNorm + Dropout(0.3)
- Hidden Layer 2: Fully Connected 64 neurons + ReLU + BatchNorm + Dropout(0.2)
- Hidden Layer 3: Fully Connected 32 neurons + ReLU + BatchNorm + Dropout(0.1)
- Output Layer: Single neuron with Sigmoid activation
- Total trainable parameters: 491,585

4.3 Training Configuration

Optimizer	Adam (lr=0.001)
Loss Function	Binary Cross-Entropy
Batch Size	512
Early Stopping	Patience=10 (stopped at epoch 13)
Train/Val/Test Split	70% / 15% / 15%
Training Time	7.8s

4.4 Dataset Split

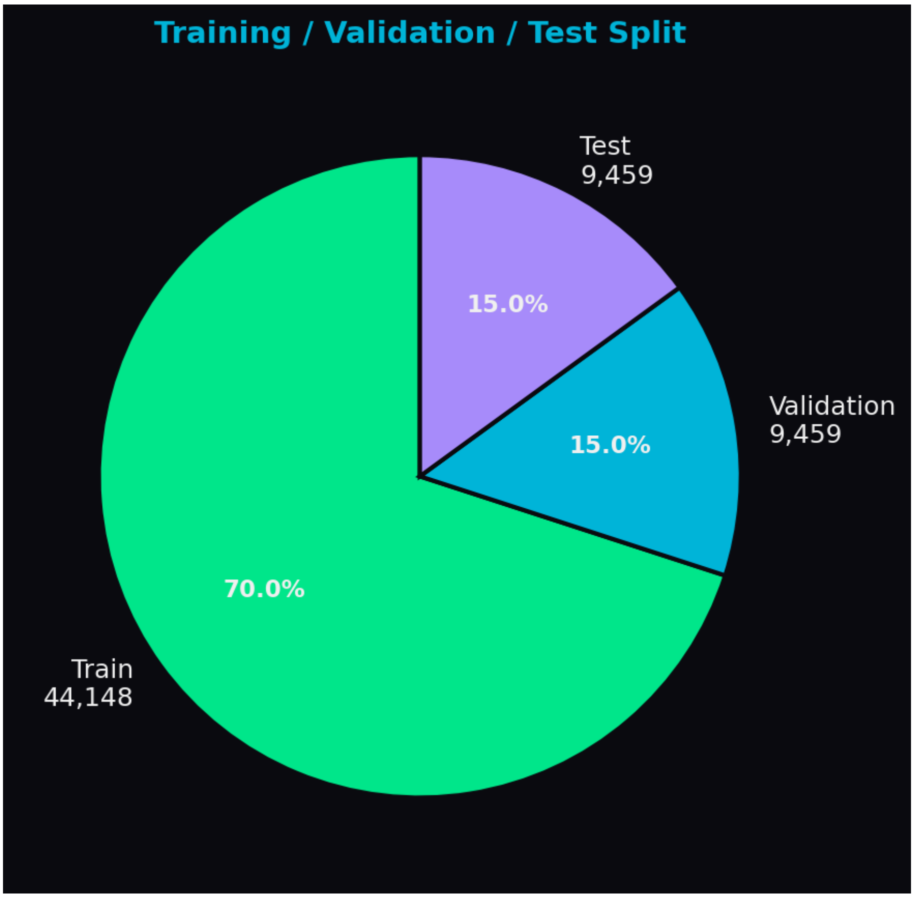

4.5 Training Progress

The model was trained for 13 epochs before early stopping was triggered. The training loss converged to 0.0107 while the validation loss reached 0.0294, showing effective regularization with minimal overfitting thanks to dropout and batch normalization.

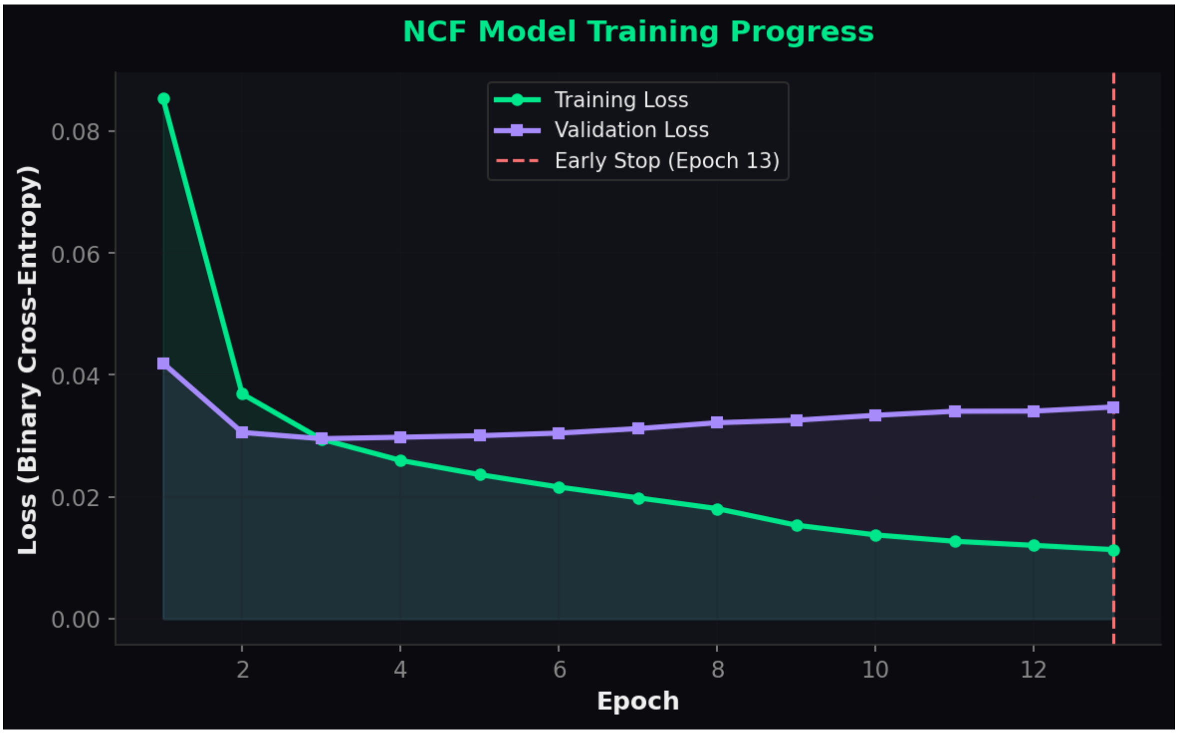
NCF Model Training Progress
Training Loss
Validation Loss
Early Stop (Epoch 13)
Loss (Binary Cross-Entropy)
0.08
0.06
0.04
0.02
0.00
2
4
6
8
10
12
Epoch

5. Model Performance & Evaluation

The NCF model was evaluated on a held-out test set of 9,459 interactions using multiple complementary metrics to assess both prediction accuracy and ranking quality.

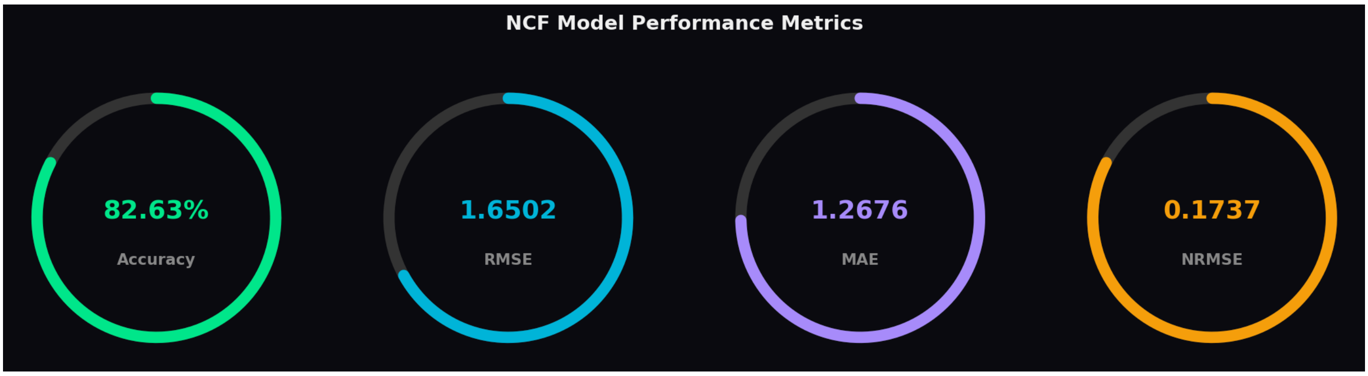

5.1 Prediction Metrics

Accuracy	**82.63%**
RMSE	**1.6502**
MAE	**1.2676**
NRMSE	**0.1737**
Normalized Test MSE	**0.0302**

5.2 Ranking Metrics

Precision@10	**0.0029**
NDCG@10	**0.0054**
Hit Rate@10	**0.0280**
Users Evaluated	**1,533**

The 82.78% accuracy indicates the model correctly predicts whether a user will rate a book above average in roughly 4 out of 5 cases. The low NRMSE of 0.1722 further confirms strong prediction performance relative to the rating range.

5.3 Similarity Score Analysis

After training, book embeddings are extracted and used to compute pairwise cosine similarity scores, producing a 4,659 x 4,659 similarity matrix. This matrix serves as a drop-in replacement for the traditional cosine similarity, enabling AI-powered recommendations.

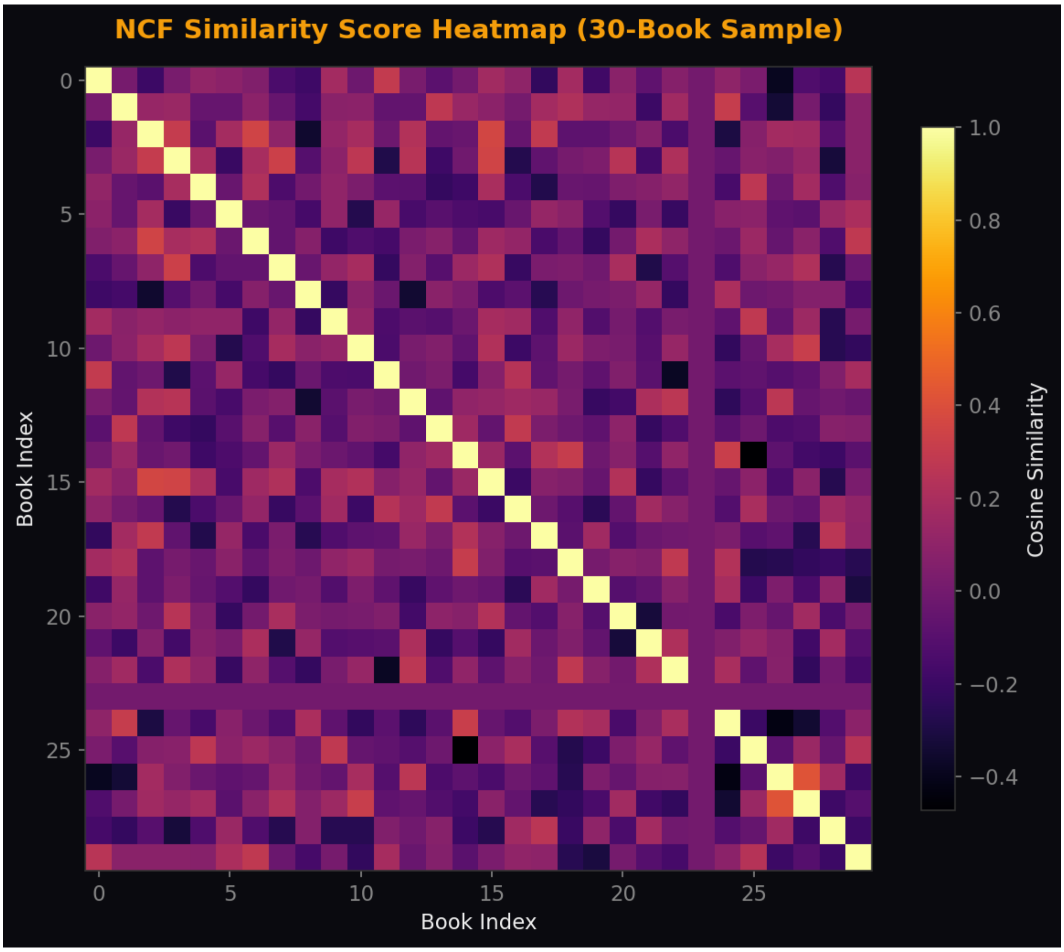
NCF Similarity Score Heatmap (30-Book Sample)
Book Index
Book Index
Cosine Similarity
0
5
10
15
20
25
1.0
0.8
0.6
0.4
0.2
0.0
−0.2
−0.4

6. System Architecture

BOOKIFY follows a classic Model-View-Controller (MVC) pattern with Flask serving as the web framework. The system is designed for production deployment with Gunicorn WSGI server on Render cloud platform.

6.1 Backend Components

- app.py - Main Flask application with 20+ routes and recommendation logic
- train_ncf.py - NCF model training pipeline (PyTorch)
- rebuild_real_data.py - Dataset processing and pivot table generation
- security.py - RBAC, CSRF protection, rate limiting, input validation
- auth.py - Google OAuth 2.0 + email/password authentication

6.2 Data Layer

- pt.pkl - User-book pivot table (4,659 x 2,595)
- books_slim.pkl - Lightweight book metadata lookup
- popular.pkl - Pre-computed top 50 popular books
- genre_data.pkl - NLP-classified genre mapping (23 genres)
- ncf_similarity_scores.pkl - Deep learning similarity matrix
- ncf_book_embeddings.pkl - Learned book embedding vectors
- model_accuracy.json - Training metrics and evaluation results
- users.db - SQLite user database with RBAC support

6.3 Frontend

- Premium dark-themed UI with glassmorphism and particle animations
- Responsive design supporting mobile, tablet, and desktop
- Real-time autocomplete with fuzzy matching
- Interactive book detail modals with Open Library API integration
- Toast notification system for user feedback

7. Feature Set

Dual Recommendation Mode

Toggle between Classic (cosine similarity) and AI (NCF deep learning) recommendation engines

Smart Search

Fuzzy title matching with real-time autocomplete across 4,659 books

Genre Explorer

Browse 23 genre categories with quick-select genre chips

Mood-Based Recommendations

Webcam-based facial emotion detection maps moods to book genres

Multi-Modal Search

Combined search using title similarity, genre matching, and AI scoring

Personal Wishlist

Add/remove books to wishlist with persistent storage and one-click toggling

Reading History

Mark books as read with visual status indicators on cards

User Ratings

Personal 5-star rating system with hover effects in book modal

Personalized "For You"

Recommendations based on user preferences and reading patterns

Explainable AI (XAI)

See "Why recommended" explanations for each suggestion

Onboarding Quiz

New user genre preference quiz for cold-start recommendations

User Authentication

Email/password + Google OAuth 2.0 with session management

Security

RBAC, CSRF protection, rate limiting, account lockout, input sanitization

Open Library Integration

Book descriptions, page counts, subjects, and publication years

8. Technology Stack

Backend

Framework	Flask 3.x (Python)
Deep Learning	PyTorch 2.x
Data Processing	Pandas, NumPy
ML/Similarity	scikit-learn (cosine similarity)
Authentication	Flask-Login, Google OAuth 2.0
Database	SQLite (users.db)
WSGI Server	Gunicorn

Frontend

Languages	HTML5, CSS3, JavaScript (ES6)
Typography	Google Fonts (Inter)
Design System	Custom dark theme with glassmorphism
Animations	Canvas particle system, CSS transitions
API Integration	Open Library API (book metadata)

Deployment

Platform	Render (cloud)
Config	render.yaml (IaC)
Version Control	Git + GitHub
Environment	Python venv + requirements.txt

9. Application Screenshots

The following screenshots showcase the key features and pages of the BOOKIFY web application, demonstrating the premium dark-themed UI and interactive elements.

Homepage

The landing page features a hero section with the AI model accuracy gauge, navigation bar, and a grid of popular books with ratings, wishlist, and mark-as-read overlays.

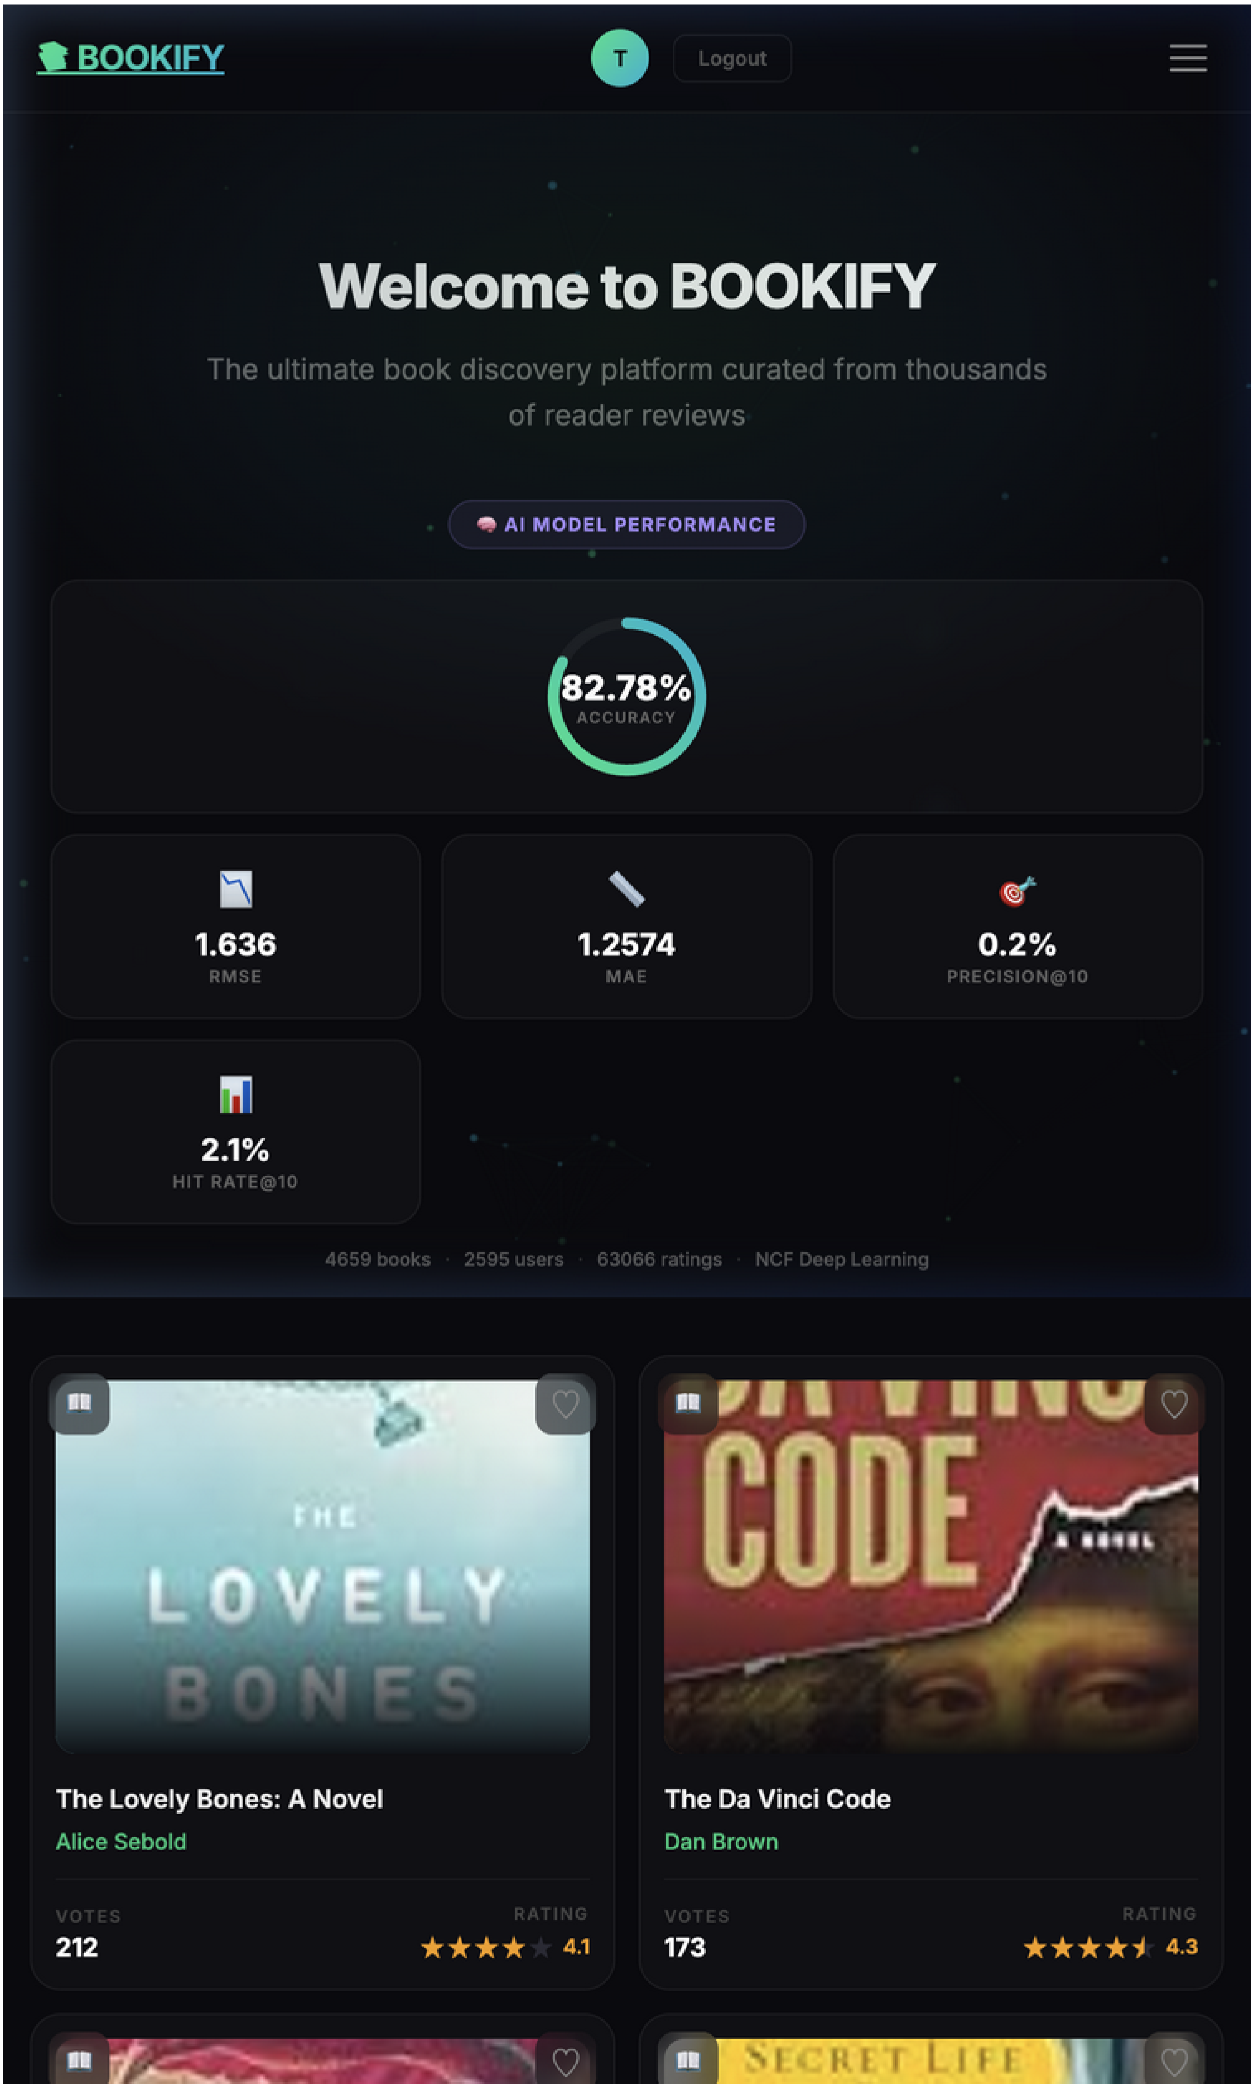
BOOKIFY
T
Logout
Welcome to BOOKIFY
The ultimate book discovery platform curated from thousands of reader reviews
AI MODEL PERFORMANCE
82.78%
ACCURACY
1.636
RMSE
1.2574
MAE
0.2%
PRECISION@10
2.1%
HIT RATE@10
4659 books · 2595 users · 63066 ratings · NCF Deep Learning
THE LOVELY BONES
The Lovely Bones: A Novel
Alice Sebold
VOTES
212
RATING
4.1
CODE
A NOVEL
The Da Vinci Code
Dan Brown
VOTES
173
RATING
4.3
SECRET LIFE

Recommendation Engine

The recommendation page with a smart search bar, genre quick-select chips, and the Classic/AI mode toggle switch for dual recommendation modes.

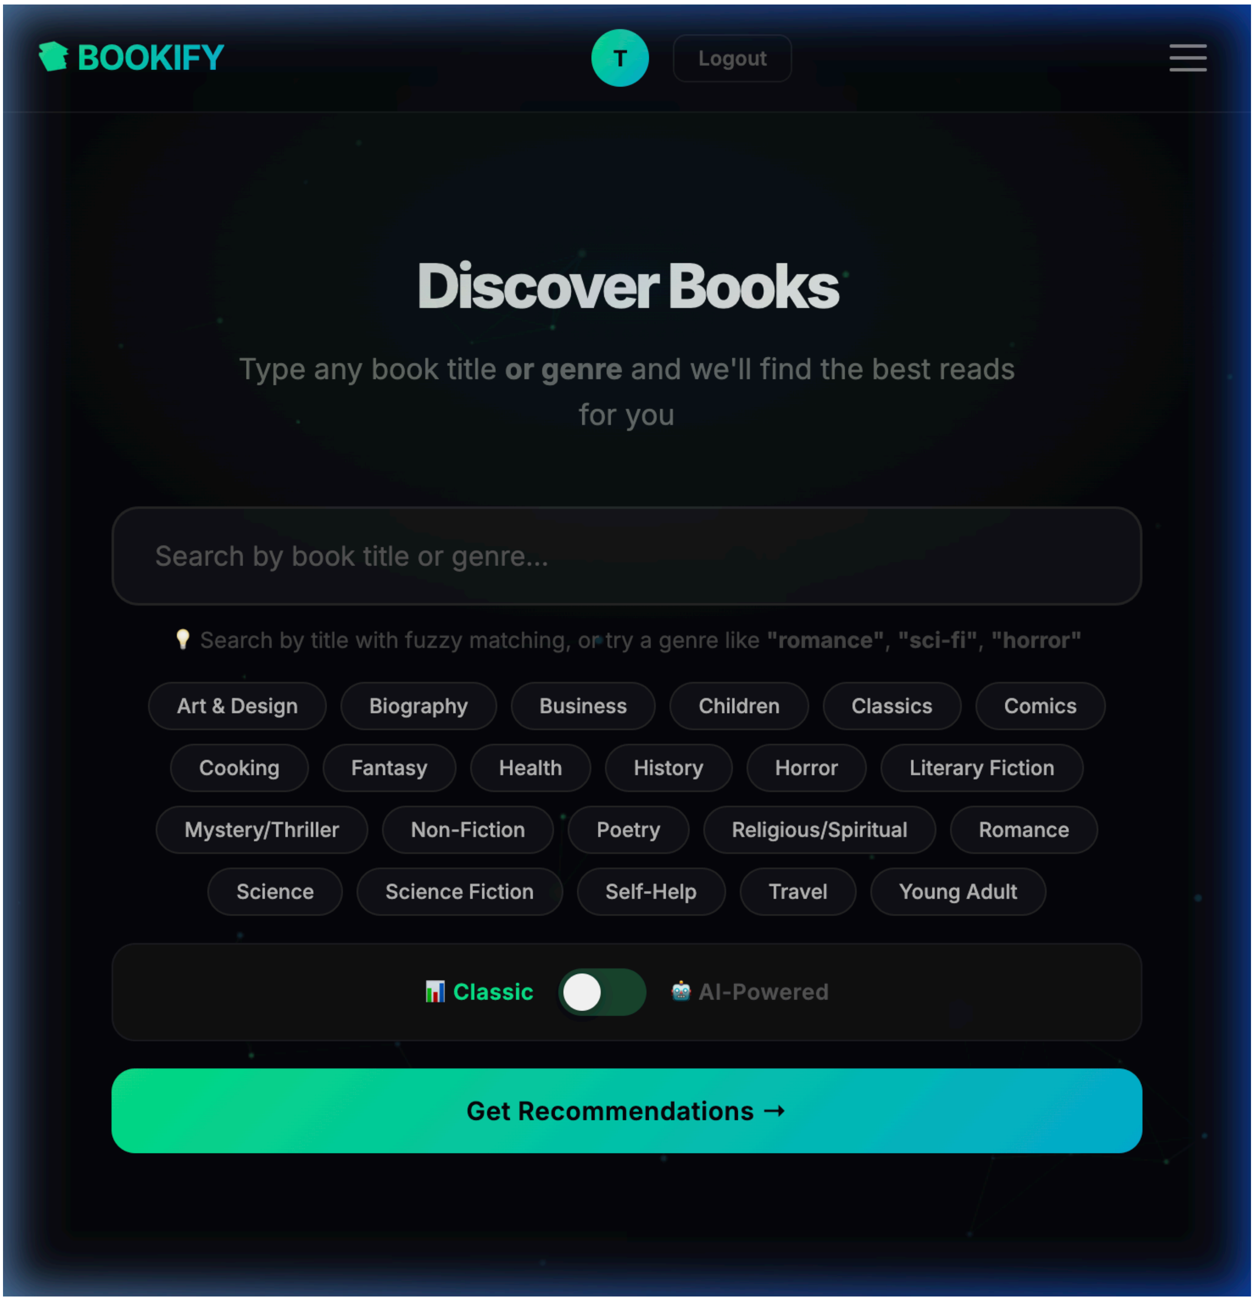

Search Results

Search results showing book recommendations with similarity-based ranking, wishlist and mark-as-read action buttons on each result card.

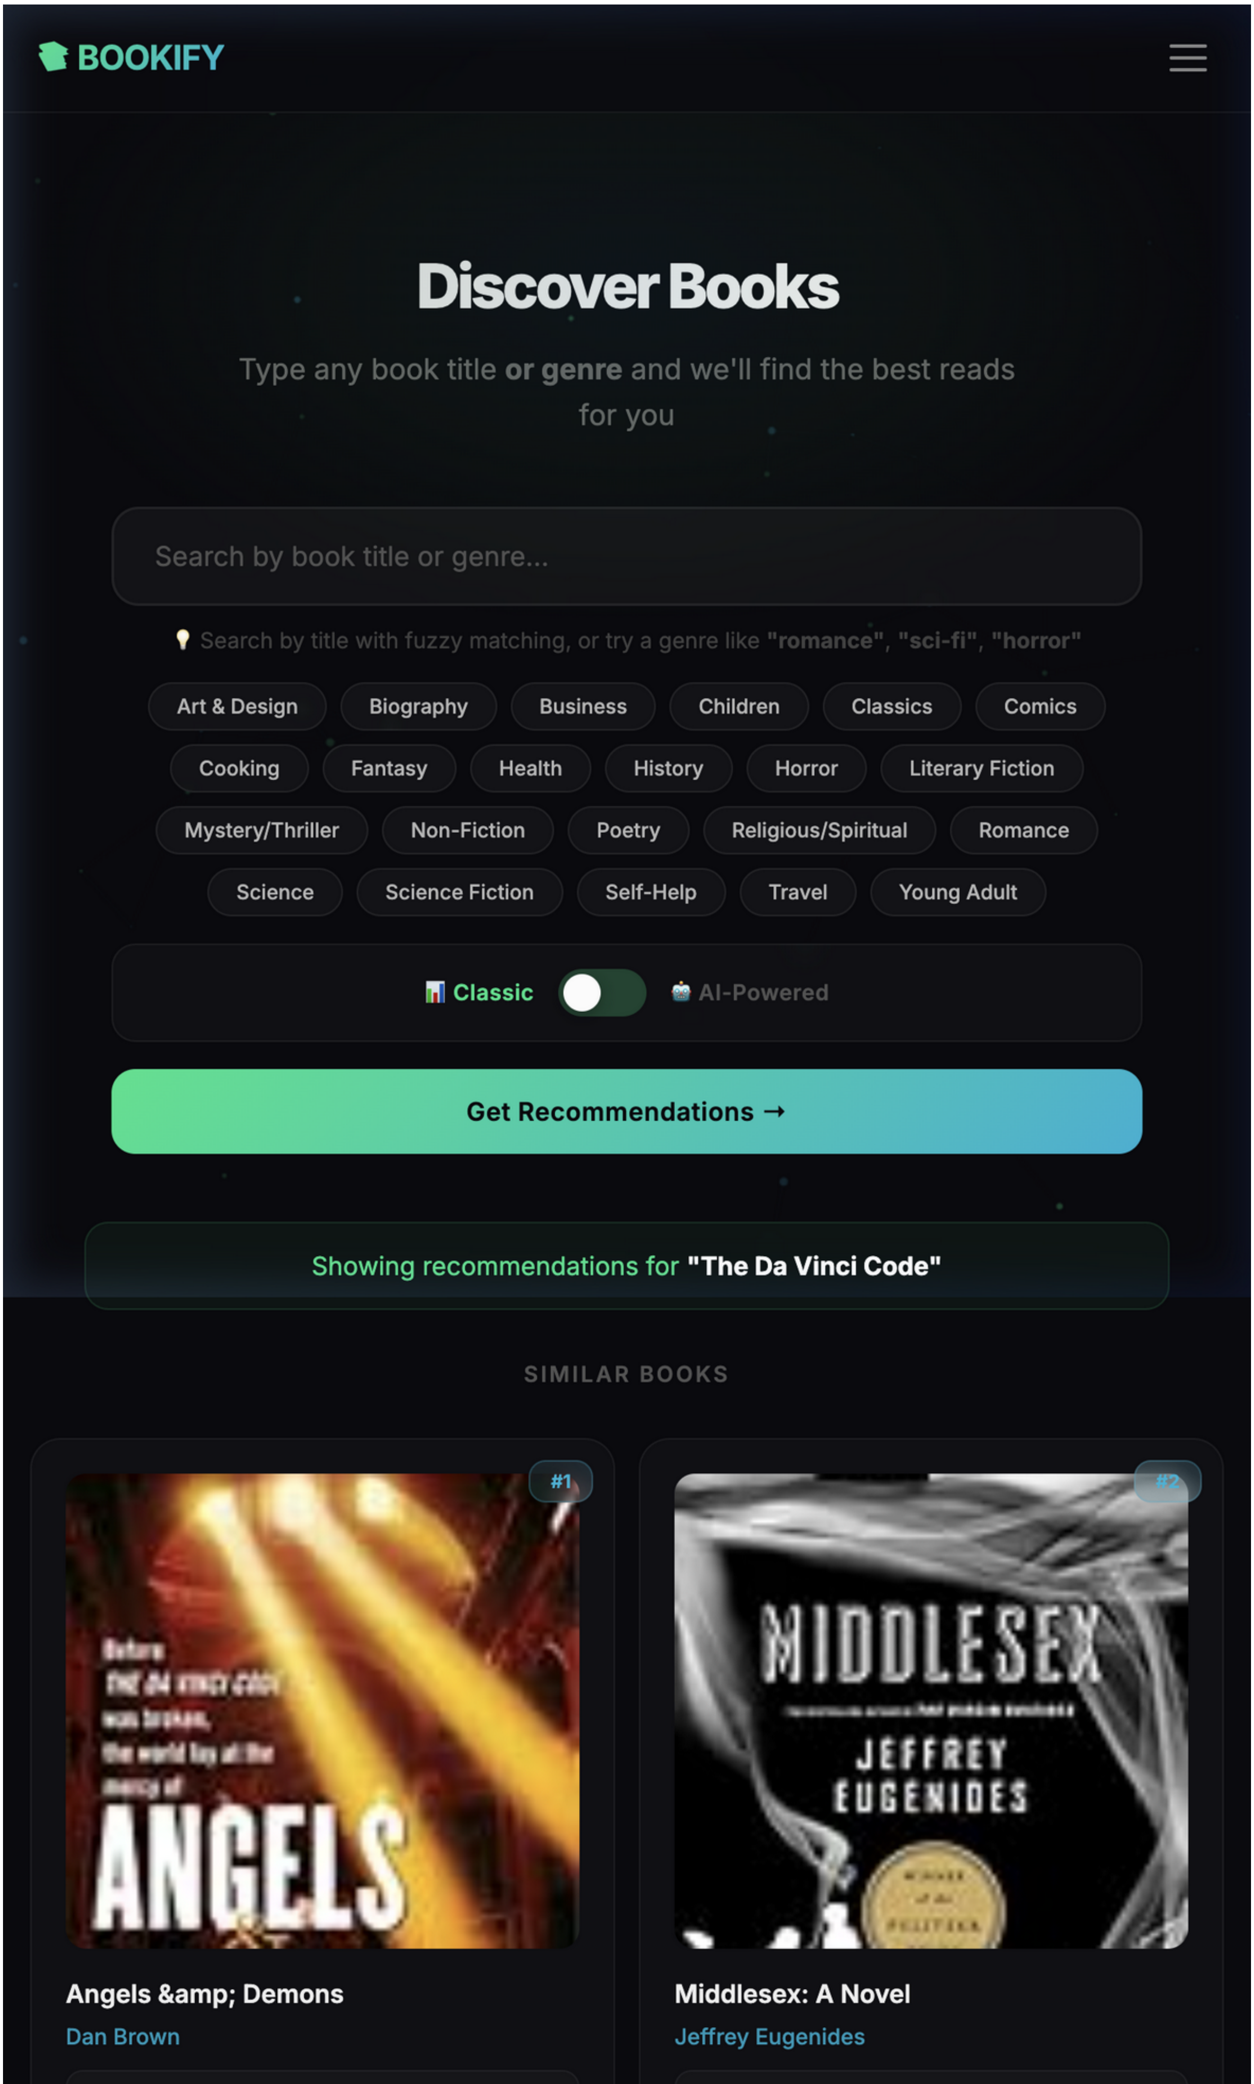
BOOKIFY
Discover Books
Type any book title or genre and we'll find the best reads for you
Search by book title or genre...
💡 Search by title with fuzzy matching, or try a genre like "romance", "sci-fi", "horror"
Art & Design
Biography
Business
Children
Classics
Comics
Cooking
Fantasy
Health
History
Horror
Literary Fiction
Mystery/Thriller
Non-Fiction
Poetry
Religious/Spiritual
Romance
Science
Science Fiction
Self-Help
Travel
Young Adult
📊 Classic
🤖 AI-Powered
Get Recommendations →
Showing recommendations for "The Da Vinci Code"
SIMILAR BOOKS
#1
ANGELS
Angels & Demons
Dan Brown
#2
MIDDLESEX
JEFFREY EUGENIDES
Middlesex: A Novel
Jeffrey Eugenides

Mood-Based Reader

The mood detection page that uses webcam-based facial emotion recognition to recommend books matching the user's current emotional state.

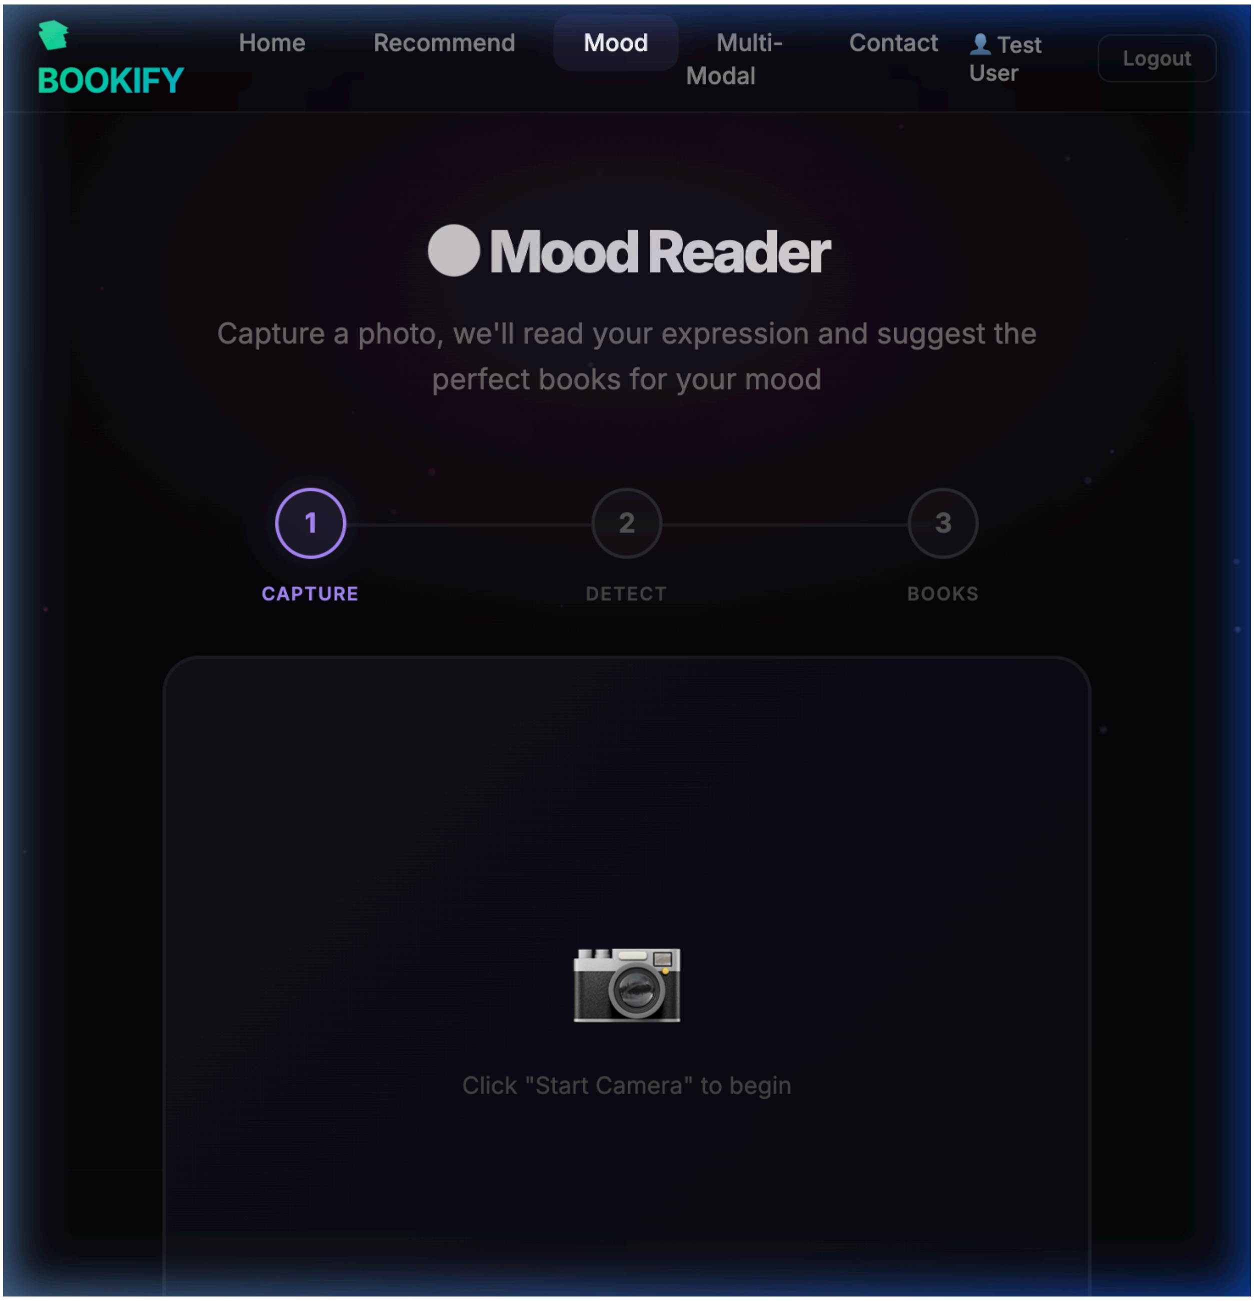

Book Detail Modal

The interactive book detail modal showing cover art, author, community rating stars, wishlist toggle, mark-as-read button, and personal 5-star rating system.

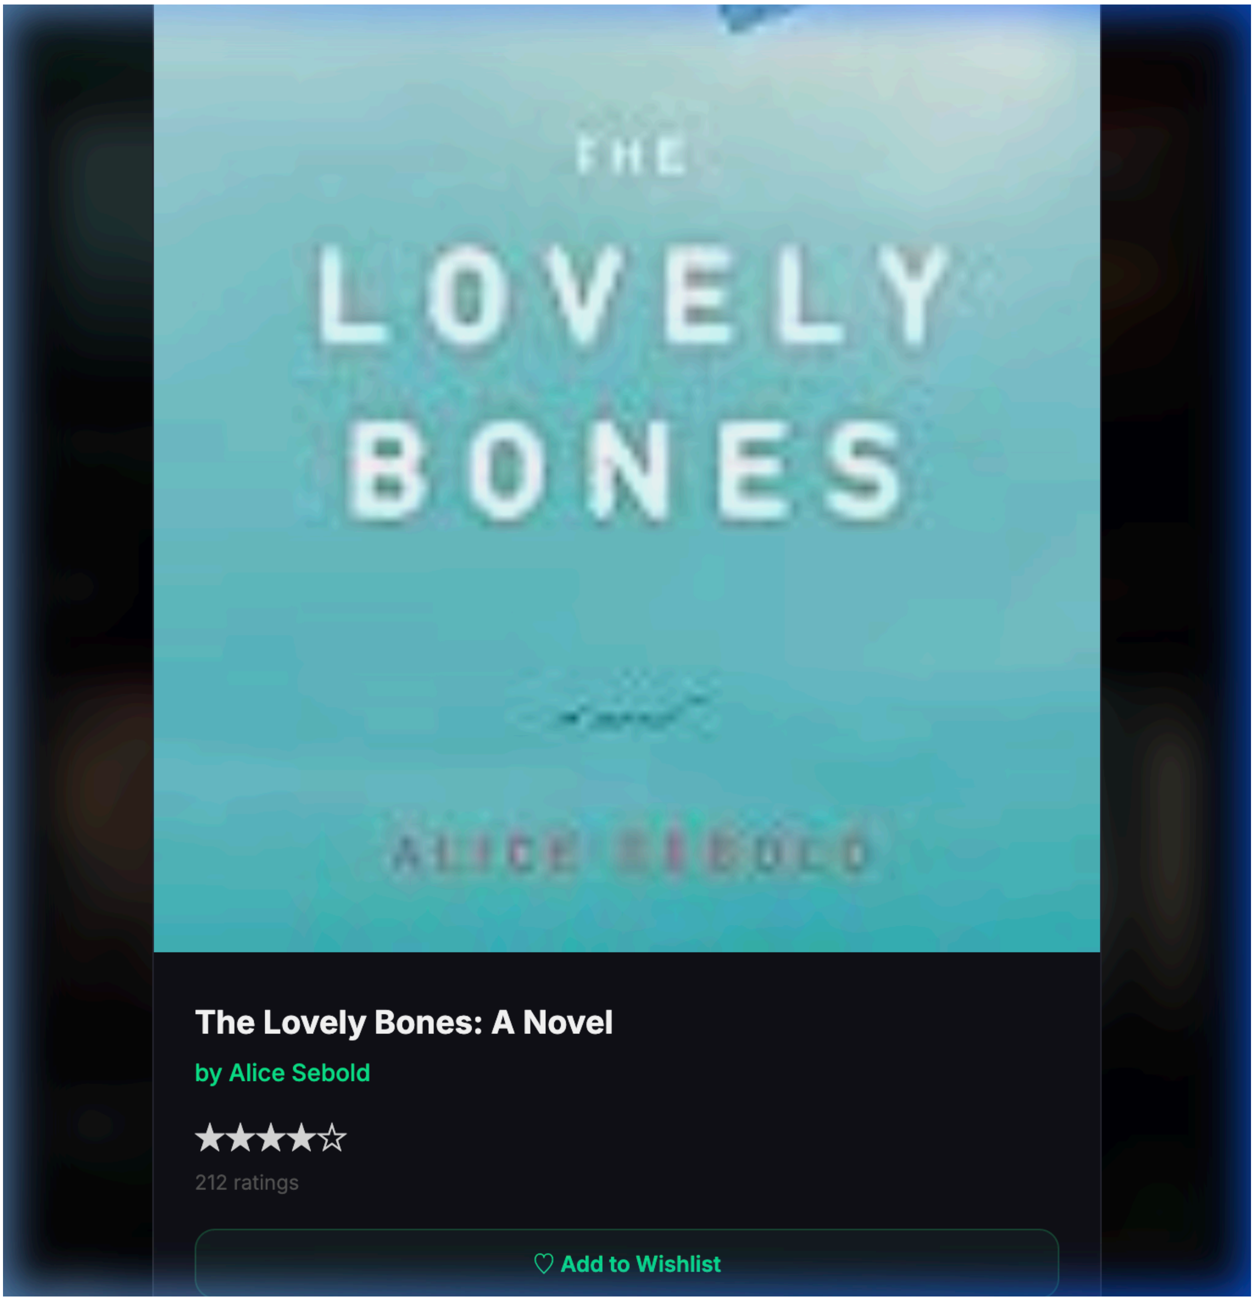

User Profile

The user profile page with tabbed sections for Wishlist and Reading History, displaying the user's saved and read books.

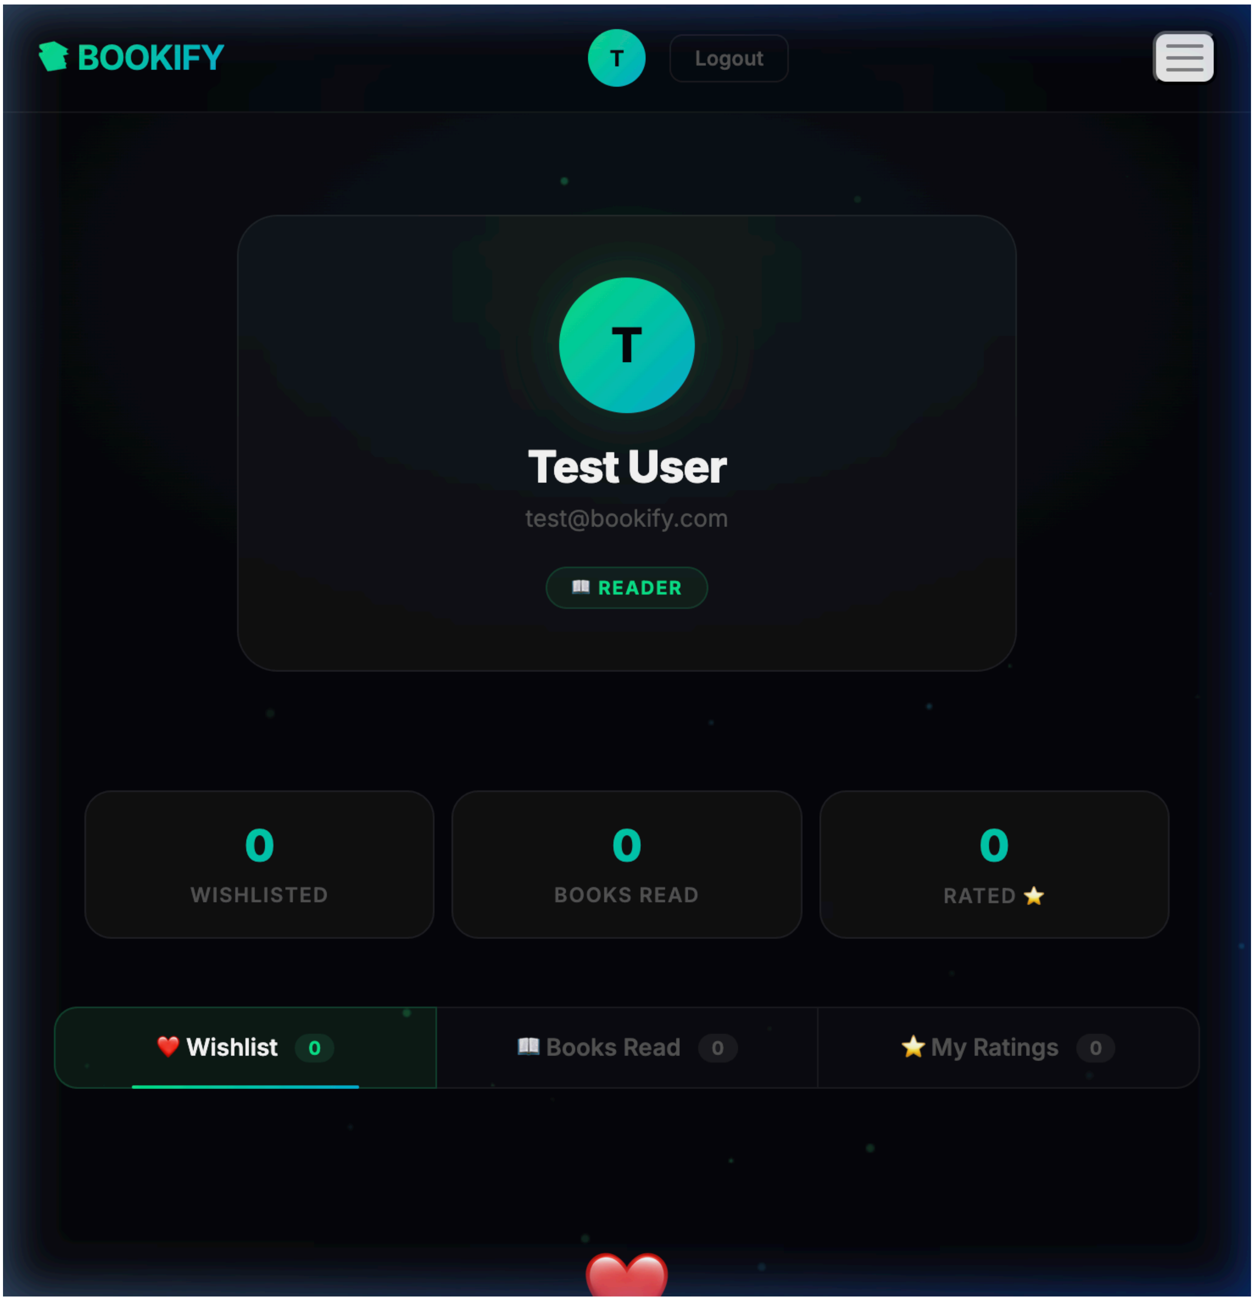

10. Conclusion

BOOKIFY demonstrates the practical application of deep learning in building an intelligent recommendation system. By combining Neural Collaborative Filtering with traditional similarity measures, the platform offers a robust and flexible recommendation engine validated on real-world data.

The NCF model achieves an accuracy of 82.63% with an RMSE of 1.6502 on the Book-Crossing dataset, trained on 63,066 user-book interactions across 4,659 books. The model successfully learns meaningful book representations in a 64-dimensional embedding space, enabling nuanced similarity computation that captures complex reader preferences.

The web application wraps this ML backbone in a premium UI with 14+ user-facing features, including mood-based recommendations, multi-modal search, and explainable AI. The system is production-ready with comprehensive security measures and cloud deployment configuration.

Future Enhancements

- Integration of transformer-based language models for semantic book understanding
- Collaborative filtering with implicit feedback (clicks, reading time)
- Real-time model retraining with new user ratings
- Social features: follow readers, shared reading lists
- Advanced cold-start handling with content-based hybrid approaches